

Observations


- Optimal BST might not have smallest height.
- Optimal BST might not have highest-probability key at root.

Build by exhaustive checking?

- Construct each n -node BST.
- For each, put in keys.
- Then compute expected search cost.
- But there are $\Omega(4^n / n^{3/2})$ different BSTs with n nodes.

Step 1: The structure of an optimal binary search tree

Consider any subtree of a BST. It contains keys in a contiguous range k_i, \dots, k_j for some $1 \leq i \leq j \leq n$.




If T is an optimal BST and T contains subtree T' with keys k_i, \dots, k_j , then T' must be an optimal BST for keys k_i, \dots, k_j .

Proof Cut and paste. ■

Use optimal substructure to construct an optimal solution to the problem from optimal solutions to subproblems:

- Given keys k_i, \dots, k_j (the problem).
- One of them, k_r , where $i \leq r \leq j$, must be the root.
- Left subtree of k_r contains k_i, \dots, k_{r-1} .
- Right subtree of k_r contains k_{r+1}, \dots, k_j .



- If
 - you examine all candidate roots k_r , for $i \leq r \leq j$, and
 - you determine all optimal BSTs containing k_i, \dots, k_{r-1} and containing k_{r+1}, \dots, k_j ,
 then you're guaranteed to find an optimal BST for k_i, \dots, k_j .