



Database Support for Data Warehouses

- Data in warehouses usually append only, not updated
 - Can avoid concurrency control overheads
- Data warehouses often use **column-oriented storage**
 - E.g. a sequence of *sales* tuples is stored as follows
 - Values of *item_id* attribute are stored as an array
 - Values of *store_id* attribute are stored as an array,
 - And so on
 - Arrays are compressed, reducing storage, IO and memory costs significantly
 - Queries can fetch only attributes that they care about, reducing IO and memory cost
 - More details in Section 13.6
- Data warehouses often use parallel storage and query processing infrastructure
 - Distributed file systems, Map-Reduce, Hive, ...