

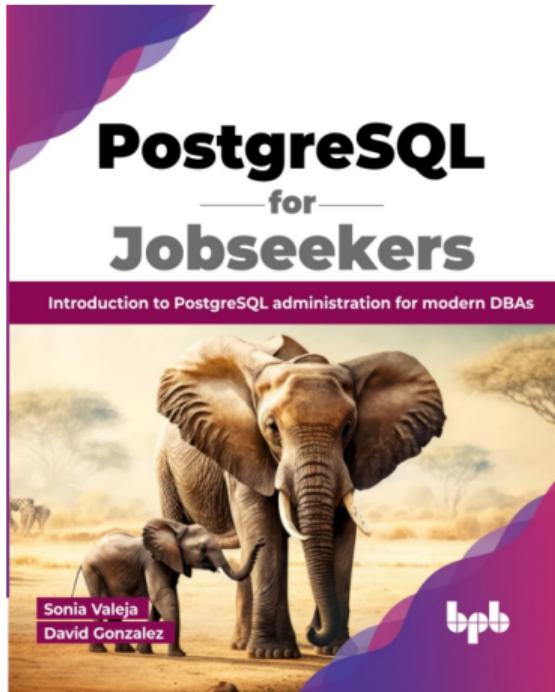
# Database Administration

## Lecture 07: Backup and Restore.

Valeja & Gonzales.

31 de agosto de 2025

# Database Administration: Backup and Restore.



Content has been extracted from *PostgreSQL for Jobseekers* (Chapter 7), by Sonia Valeja and David Gonzales, 2023. Visit <https://bpbonline.com/products/postgresql-for-jobseekers>.

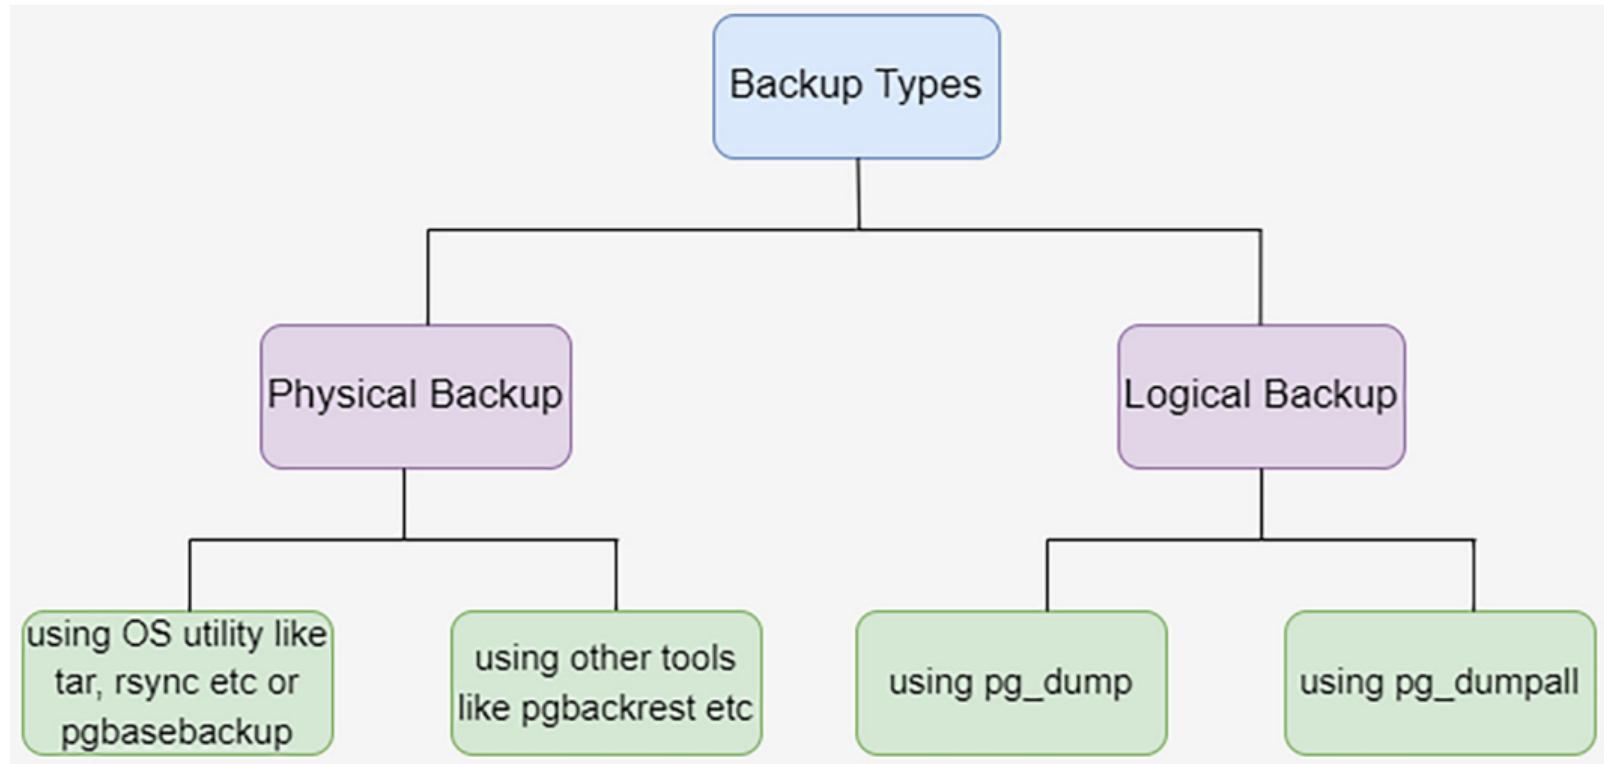
# Introduction

- ▶ Regular backups are crucial for database administration.
- ▶ Backups enable recovery from accidental changes.
- ▶ PostgreSQL provides multiple backup and restore tools.

# Types of Backups

- ▶ **Physical Backup:** Copies entire data directory.
- ▶ **Logical Backup:** Exports data in human-readable format.

# Types of Backups



# Physical Backup

- ▶ Uses filesystem-level copying.
- ▶ Steps:
  1. Create a checkpoint.
  2. Copy data directory (e.g., using `tar` or `rsync`).
  3. Stop backup process.
- ▶ Example tool: `pg_basebackup`.

## pg\_basebackup

```
postgres@ip-172-31-27-8:~$ pg_basebackup --help
```

`pg_basebackup` takes a base backup of a running PostgreSQL server.

## Usage:

`pg_basebackup [OPTION] ...`

## Options controlling the output:

**-D, --pgdata=DIRECTORY** receive base backup into directory

**-F, --format=plt** output format (plain (default), tar)

**-r, --max-rate=RATE** maximum transfer rate to transfer data directory  
(in kB/s, or use suffix "k" or "M")

**-R, --write-recovery-conf** write configuration for replication

**-T, --tablespace-mapping=OLDDIR=NEWDIR**  
relocate tablespace in OLDDTR to NEWDTR

`--wal-dir=WALDIR` location for the write-ahead log directory

**-X, --wal-method=none|fetch|stream** include required WAL files with specified method

# Point in Time Recovery (PITR) / Archival

**PITR** by enabling **archiving**, this allows restoring the database to a specific time frame, provided that WAL<sup>1</sup>/archives are restored until that point.

## Steps:

- ▶ Enable archiving in `postgresql.conf`.
- ▶ Configure the following parameters:

## Configuration Example:

```
archive_mode = on
archive_command = 'test ! -f /mnt/server/archivedir/%f && \
                  cp %p /mnt/server/archivedir/%f' # Linux
archive_command = 'copy "%p" "C:\\server\\archivedir\\%f"' # Windows
```

We will see in the restore section how to recover a database to a specific time frame.

---

<sup>1</sup>Write-Ahead Logging

# Pros and Cons of Physical Backup

## Pros:

- ▶ Contains a full filesystem-level backup of the data directory.
- ▶ Includes configuration files such as `postgresql.conf`, `pg_hba.conf`, etc.
- ▶ Simple process: copy the data directory and custom tablespaces to the destination server and start PostgreSQL.

## Cons:

- ▶ Only cluster-level backups are possible (no single database or table backup).
- ▶ Consistency of the backup is difficult to guarantee.
- ▶ Backup/restore speed depends heavily on server hardware and network speed (for remote copies).

# Logical Backup

- ▶ Exports data in SQL format.
- ▶ Allows individual table or database backups.
- ▶ Example tools:
  - ▶ pg\_dump for single database.
  - ▶ pg\_dumpall for full cluster backup.

# pg\_dump

```
postgres@ip-172-31-27-8:~$ pg_dump --help
pg_dump dumps a database as a text file or to other formats.
```

Usage:

```
pg_dump [OPTION]... [DBNAME]
```

General options:

-f, --file=FILENAME	output file or directory name
-F, --format=c\ldlt\p	output file format (custom, directory, tar, plain text (default))
-j, --jobs=NUM	use this many parallel jobs to dump
-v, --verbose	verbose mode
-V, --version	output version information, then exit
-Z, --compress=0-9	compression level for compressed formats
--lock-wait-timeout=TIMEOUT	fail after waiting TIMEOUT for a table lock
--no-sync	do not wait for changes to be written safely to disk
-?, --help	show this help, then exit

Options controlling the output content:

-a, --data-only	dump only the data, not the schema
-b, --blobs	include large objects in dump
-B, --no-blobs	exclude large objects in dump
-c, --clean	clean (drop) database objects before recreating

# Pros and Cons of Logical Backup

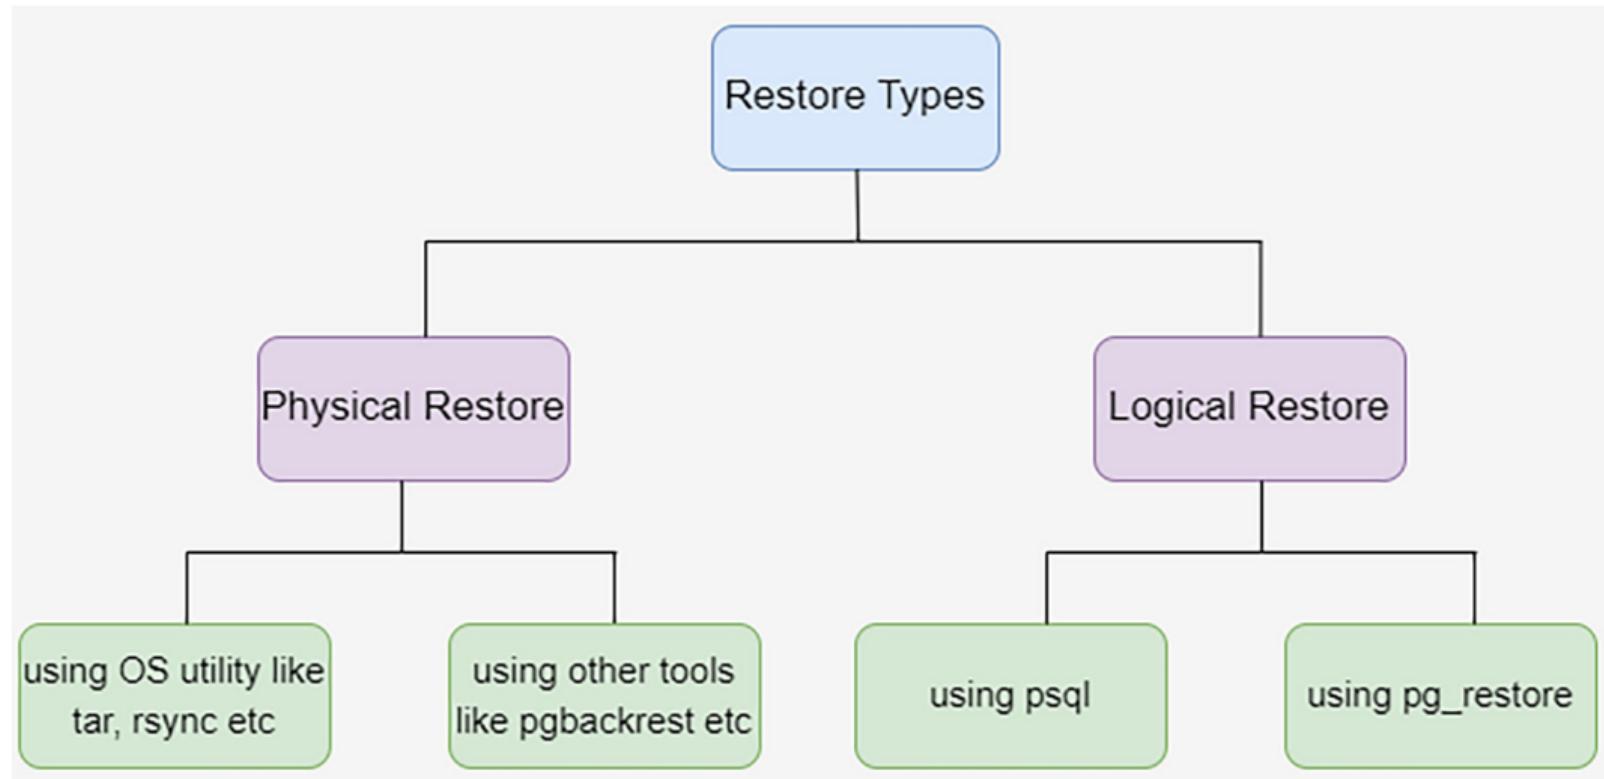
## Pros:

- ▶ Can back up specific parts of the database (not just the entire cluster).
- ▶ Simple and human-readable format.
- ▶ Backups are stored as plain text/SQL files, making them portable across different operating systems.
- ▶ Faster than physical backup when only a single DB object is required.

## Cons:

- ▶ Configuration files are not included (must be backed up separately).
- ▶ Point-in-Time Recovery (PITR) is not supported.

# Types of Restore



## Restore Methods

- ▶ `psql`: Restores backups in plain SQL format.
- ▶ `pg_restore`: Restores backups from compressed/custom format.

# pg\_restore

```
postgres@ip-172-31-27-8:~$ pg_restore --help
pg_restore restores a PostgreSQL database from an archive created by pg_dump.
```

Usage:

```
pg_restore [OPTION]... [FILE]
```

General options:

-d, --dbname=NAME	connect to database name
-f, --file=FILENAME	output file name (- for stdout)
-F, --format=cldlt	backup file format (should be automatic)
-l, --list	print summarized TOC of the archive
-v, --verbose	verbose mode
-V, --version	output version information, then exit
-?, --help	show this help, then exit

Options controlling the restore:

-a, --data-only	restore only the data, no schema
-c, --clean	clean (drop) database objects before recreating
-C, --create	create the target database
-e, --exit-on-error	exit on error, default is to continue
-I, --index=NAME	restore named index

# Logical Backup and Restore example

Description	Command
To dump a database called mydb into a SQL-script file:	<code>pg_dump mydb &gt; "/opt/mydb_dump.sql"</code>
To reload such a script into a (freshly created) database named newdb:	<code>psql -d newdb -f "/opt/mydb_dump.sql"</code>
To dump a database into a custom-format archive file:	<code>pg_dump -Fc mydb &gt; "/opt/mydb_dump.dump"</code>
To reload an archive file into a (freshly created) database named newdb:	<code>pg_restore -d newdb "/opt/mydb_dump.dump"</code>
To dump a single table named mytab	<code>pg_dump -t mytab mydb &gt; "/opt/mydb_dump.sql"</code>
To dump the database with the log file using user test	<code>pg_dump -U test -v -f "/opt/mydb_dump.sql" mydb 2&gt;&gt; "/opt/mydb_log.log"</code> password for user test :
To dump only structure without data	<code>pg_dump -sU test -v -f "/opt/mydb_dump.sql" mydb 2&gt;&gt; "/opt/mydb_log.log"</code> password for user test :
To take insert scripts of particular table	<code>pg_dump --column-inserts -a -t mytab -U test mydb &gt; "/opt/mytab_inserts.sql"</code>
To dump only specific tables with data	<code>pg_dump -U test -n schema1 -t BCL_* -f "/opt/BCL_TABLES.sql" mydb</code>

# Useful Backup and Restore Tools

## Community-developed tools:

- ▶ Support full, incremental, and differential backups.
- ▶ Provide backup compression and checksum.
- ▶ Allow multiple repositories: on-premise or cloud.
- ▶ Enable parallelism for faster operations.
- ▶ Maintain a backup catalog.
- ▶ Offer alerting and notification features.

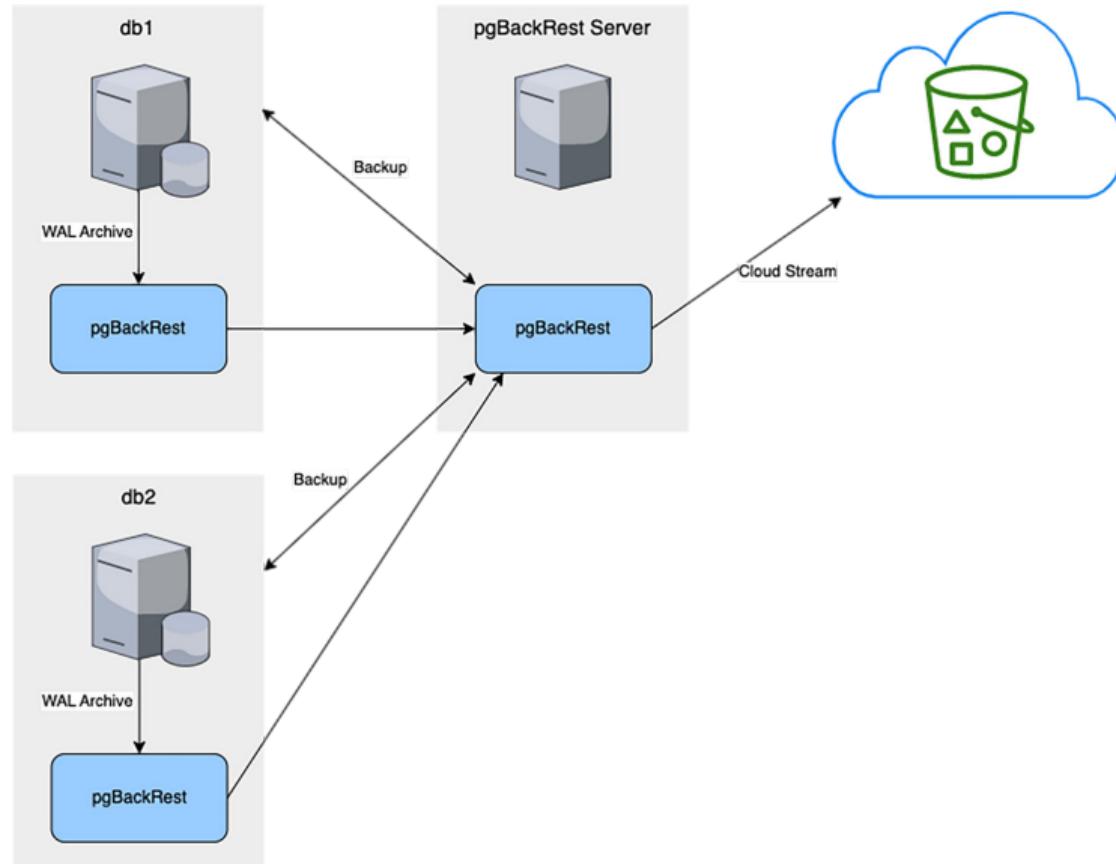
*We will next review some of the most popular open-source tools, their options, and how to perform a backup and restore.*

# Advanced Backup Tools

## pgBackRest

- ▶ Supports full, incremental, and differential backups.
- ▶ Uses parallelism for faster operations.
- ▶ Supports cloud storage.

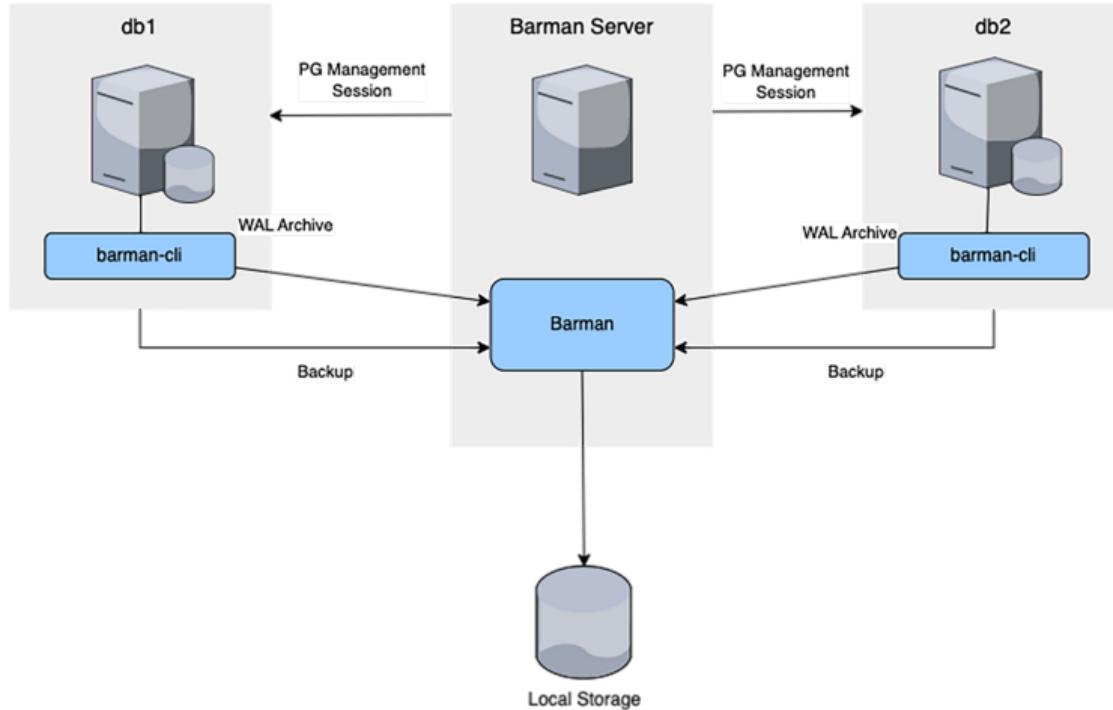
# pgBackRest



## Barman

- ▶ Used for remote backups.
- ▶ Provides incremental backup and deduplication.
- ▶ Maintains a backup catalog.

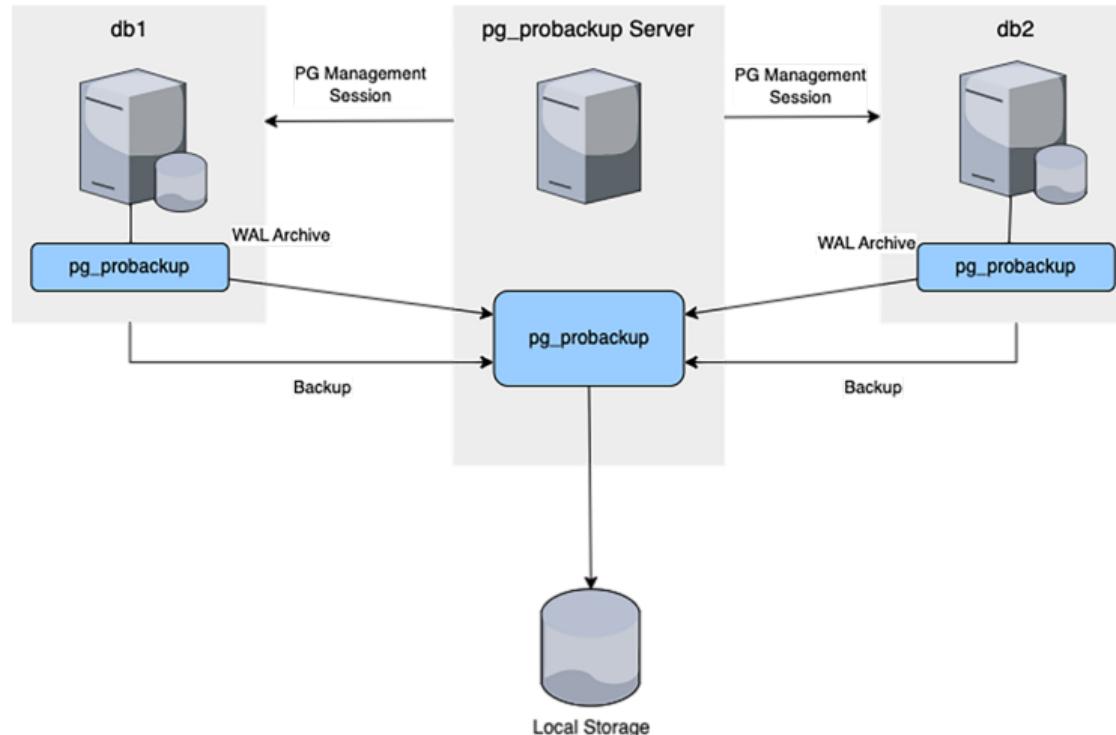
# Barman



## pg\_probackup

- ▶ Provides incremental restore to save time.
- ▶ Supports backup merging and deduplication.
- ▶ Offers parallelism and remote operations.

# pg\_probackup



TDT5FTOTC



TDT5FTOTC

- 5 **Regular backups** are essential for ensuring high availability, enabling point-in-time recovery, and recovering from accidental changes.

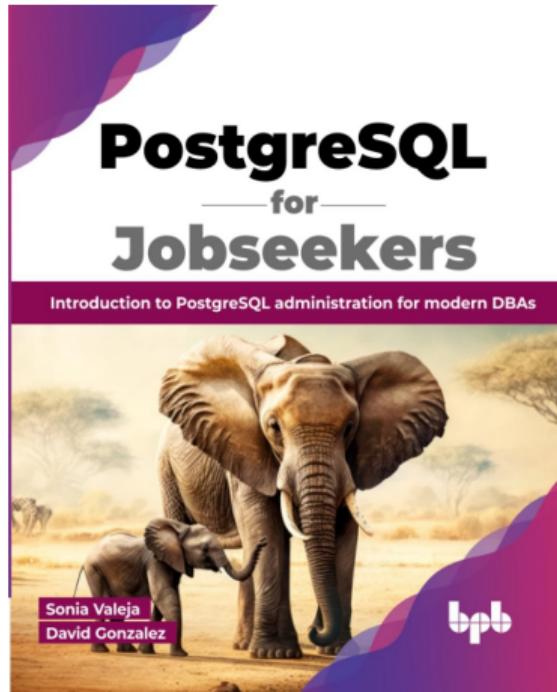
- 5 **Regular backups** are essential for ensuring high availability, enabling point-in-time recovery, and recovering from accidental changes.
- 4 PostgreSQL supports **two main backup types**: physical (cluster-level, filesystem copy) and logical (human-readable, object-level, portable).

- 5 **Regular backups** are essential for ensuring high availability, enabling point-in-time recovery, and recovering from accidental changes.
- 4 PostgreSQL supports **two main backup types**: physical (cluster-level, filesystem copy) and logical (human-readable, object-level, portable).
- 3 **Built-in tools** like pg\_dump, pg\_dumpall, pg\_basebackup, psql, and pg\_restore allow creating and restoring different types of backups.

- 5 **Regular backups** are essential for ensuring high availability, enabling point-in-time recovery, and recovering from accidental changes.
- 4 PostgreSQL supports **two main backup types**: physical (cluster-level, filesystem copy) and logical (human-readable, object-level, portable).
- 3 **Built-in tools** like pg\_dump, pg\_dumpall, pg\_basebackup, psql, and pg\_restore allow creating and restoring different types of backups.
- 2 **Point-in-Time Recovery (PITR)** requires WAL archiving and is only supported with physical backups, not logical ones.

- 5 **Regular backups** are essential for ensuring high availability, enabling point-in-time recovery, and recovering from accidental changes.
- 4 PostgreSQL supports **two main backup types**: physical (cluster-level, filesystem copy) and logical (human-readable, object-level, portable).
- 3 **Built-in tools** like pg\_dump, pg\_dumpall, pg\_basebackup, psql, and pg\_restore allow creating and restoring different types of backups.
- 2 **Point-in-Time Recovery (PITR)** requires WAL archiving and is only supported with physical backups, not logical ones.
- 1 **Community tools** such as pgBackRest, Barman, and pg\_probackup provide advanced features like incremental backups, compression, cloud storage, and catalogs.

# Database Administration: Backup and Restore.



Content has been extracted from *PostgreSQL for Jobseekers* (Chapter 7), by Sonia Valeja and David Gonzales, 2023. Visit <https://bpbonline.com/products/postgresql-for-jobseekers>.