


ML101
Lecture 03: Linear Regression
An Introduction to Statistical Learning

James, Witten, Hastie & Tibshirani

13 de agosto de 2026

An Introduction to Statistical Learning with Applications in R

a.k.a. ISLR



Content has been extracted from *An Introduction to Statistical Learning with Applications in R*, Second Edition, by James, Witten, Hastie & Tibshirani. Springer. 2023.

Visit <https://www.statlearning.com/>.

Plan

Simple Linear Regression

- Estimating the Coefficients

- Assessing the Accuracy of the Coefficient Estimates

- Assessing the Accuracy of the Model

Multiple Linear Regression

- Estimating the Regression Coefficients

- Some Important Questions

Other Considerations in the Regression Model


- Qualitative Predictors

- Extensions of the Linear Model

- Potential Problems


Linear Regression

- ▶ Linear regression is a simple approach to supervised learning. It assumes that the dependence of Y on $X_1, X_2, \dots, X_p$ is linear.
- ▶ True regression functions are never linear!



Linear Regression

- ▶ Linear regression is a simple approach to supervised learning. It assumes that the dependence of Y on $X_1, X_2, \dots, X_p$ is linear.
- ▶ True regression functions are never linear!






- ▶ Although it may seem overly simplistic, linear regression is extremely useful both conceptually and practically.

Linear Regression for the Advertising Data

Consider the advertising data. Questions we might ask:

- ▶ Is there a relationship between advertising budget and sales?
- ▶ How strong is the relationship between advertising budget and sales?
- ▶ Which media contribute to sales?
- ▶ How accurately can we predict future sales?
- ▶ Is the relationship linear?
- ▶ Is there synergy among the advertising media?

Advertising data



Simple Linear Regression

Simple linear regression using a single predictor X :

- ▶ We assume a model:

$$Y = \beta_0 + \beta_1 X + \epsilon$$

where β_0 and β_1 are two unknown constants that represent the intercept and slope, also known as coefficients or parameters, and ϵ is the error term.

- ▶ Given some estimates $\hat{\beta}_0$ and $\hat{\beta}_1$ for the model coefficients, we predict future sales using:

$$\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x$$

where $\hat{y}$ indicates a prediction of Y on the basis of $X = x$. The hat symbol denotes an estimated value.

Estimation of the Parameters by Least Squares

- ▶ Let $\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$ be the prediction for Y based on the i -th value of X . Then $e_i = y_i - \hat{y}_i$ represents the i -th residual.
- ▶ We define the residual sum of squares (RSS) as:

$$RSS = e_1^2 + e_2^2 + \cdots + e_n^2$$

or equivalently as:


$$RSS = (y_1 - \hat{\beta}_0 - \hat{\beta}_1 x_1)^2 + (y_2 - \hat{\beta}_0 - \hat{\beta}_1 x_2)^2 + \cdots + (y_n - \hat{\beta}_0 - \hat{\beta}_1 x_n)^2$$

- ▶ The least squares approach chooses $\hat{\beta}_0$ and $\hat{\beta}_1$ to minimize the RSS. The minimizing values are:

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2}, \quad \hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$$

where $\bar{y} \equiv \frac{1}{n} \sum_{i=1}^n y_i$ and $\bar{x} \equiv \frac{1}{n} \sum_{i=1}^n x_i$ are the sample means.

Example: advertising data



The least squares fit for the regression of **sales** and **TV**. In this case a linear fit captures the essence of the relationship, although it is somewhat deficient in the left of the plot.

Assessing the Accuracy of the Coefficient Estimates

- The standard error of an estimator reflects how it varies under repeated sampling:

$$SE(\hat{\beta}_1)^2 = \frac{\sigma^2}{\sum_{i=1}^n (x_i - \bar{x})^2}, \quad SE(\hat{\beta}_0)^2 = \sigma^2 \left[\frac{1}{n} + \frac{\bar{x}^2}{\sum_{i=1}^n (x_i - \bar{x})^2} \right]$$

where $\sigma^2 = \text{Var}(\epsilon)$.

Assessing the Accuracy of the Coefficient Estimates

- ▶ The standard error of an estimator reflects how it varies under repeated sampling:

$$SE(\hat{\beta}_1)^2 = \frac{\sigma^2}{\sum_{i=1}^n (x_i - \bar{x})^2}, \quad SE(\hat{\beta}_0)^2 = \sigma^2 \left[\frac{1}{n} + \frac{\bar{x}^2}{\sum_{i=1}^n (x_i - \bar{x})^2} \right]$$

where $\sigma^2 = \text{Var}(\epsilon)$.

- ▶ These standard errors can be used to compute **confidence intervals**. A 95 % confidence interval is defined as a range of values such that with 95 % probability, the range will contain the true unknown value of the parameter. It has the form:

$$\hat{\beta}_1 \pm 2 \cdot SE(\hat{\beta}_1)$$

Confidence Intervals Continued

- That is, there is approximately a 95 % chance that the interval:

$$[\hat{\beta}_1 - 2 \cdot SE(\hat{\beta}_1), \hat{\beta}_1 + 2 \cdot SE(\hat{\beta}_1)]$$

will contain the true value of β_1 (under a scenario where we got repeated samples like the present sample).

Confidence Intervals Continued

- ▶ That is, there is approximately a 95 % chance that the interval:

$$[\hat{\beta}_1 - 2 \cdot SE(\hat{\beta}_1), \hat{\beta}_1 + 2 \cdot SE(\hat{\beta}_1)]$$

will contain the true value of β_1 (under a scenario where we got repeated samples like the present sample).

- ▶ For the advertising data, the 95 % confidence interval for β_1 is $[0,042, 0,053]$.

Hypothesis Testing

- Standard errors can also be used to perform **hypothesis tests** on the coefficients. The most common hypothesis test involves testing the **null hypothesis** of:

H_0 : There is no relationship between X and Y
versus the **alternative hypothesis**

H_A : There is some relationship between X and Y .

Hypothesis Testing

- Standard errors can also be used to perform **hypothesis tests** on the coefficients. The most common hypothesis test involves testing the **null hypothesis** of:

H_0 : There is no relationship between X and Y
versus the **alternative hypothesis**

H_A : There is some relationship between X and Y .

- Mathematically, this corresponds to testing:

$$H_0 : \beta_1 = 0 \quad \text{versus} \quad H_A : \beta_1 \neq 0$$

since if $\beta_1 = 0$ then the model reduces to $Y = \beta_0 + \epsilon$, and X is not associated with Y .

Hypothesis Testing – continued

- ▶ To test the null hypothesis, we compute a ***t*-statistic**, given by:

$$t = \frac{\hat{\beta}_1 - 0}{SE(\hat{\beta}_1)}$$

- ▶ This will have a *t*-distribution with $n - 2$ degrees of freedom, assuming $\beta_1 = 0$.
- ▶ Using statistical software, it is easy to compute the probability of observing any value equal to $|t|$ or larger. We call this probability the ***p*-value**.

Results for the Advertising Data

	Coefficient	Std. Error	<i>t</i>-statistic	<i>p</i>-value
Intercept	7.0325	0.4578	15.36	< 0.0001
TV	0.0475	0.0027	17.67	< 0.0001

Assessing the Overall Accuracy of the Model

- We compute the **Residual Standard Error**:

$$RSE = \sqrt{\frac{1}{n-2}RSS} = \sqrt{\frac{1}{n-2} \sum_{i=1}^n (y_i - \hat{y}_i)^2},$$

where the **residual sum-of-squares** is $RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2$.

Assessing the Overall Accuracy of the Model

- We compute the **Residual Standard Error**:

$$RSE = \sqrt{\frac{1}{n-2}RSS} = \sqrt{\frac{1}{n-2} \sum_{i=1}^n (y_i - \hat{y}_i)^2},$$

where the **residual sum-of-squares** is $RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2$.

- **R-squared** or fraction of variance explained is:

$$R^2 = \frac{TSS - RSS}{TSS} = 1 - \frac{RSS}{TSS}$$

where $TSS = \sum_{i=1}^n (y_i - \bar{y})^2$ is the **total sum of squares**.

Assessing the Overall Accuracy of the Model

- We compute the **Residual Standard Error**:

$$RSE = \sqrt{\frac{1}{n-2}RSS} = \sqrt{\frac{1}{n-2} \sum_{i=1}^n (y_i - \hat{y}_i)^2},$$

where the **residual sum-of-squares** is $RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2$.

- **R-squared** or fraction of variance explained is:

$$R^2 = \frac{TSS - RSS}{TSS} = 1 - \frac{RSS}{TSS}$$

where $TSS = \sum_{i=1}^n (y_i - \bar{y})^2$ is the **total sum of squares**.

- In simple linear regression, $R^2 = r^2$, where r is the sample correlation between X and Y :

$$r = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^n (x_i - \bar{x})^2} \sqrt{\sum_{i=1}^n (y_i - \bar{y})^2}}$$

Advertising Data Results

Quantity	Value
Residual Standard Error	3.26
R^2	0.612
F -statistic	312.1

Plan

Simple Linear Regression

- Estimating the Coefficients

- Assessing the Accuracy of the Coefficient Estimates

- Assessing the Accuracy of the Model

Multiple Linear Regression

- Estimating the Regression Coefficients

- Some Important Questions

Other Considerations in the Regression Model

- Qualitative Predictors

- Extensions of the Linear Model

- Potential Problems

Multiple Linear Regression

- ▶ Here our model is:

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_p X_p + \epsilon$$

- ▶ We interpret β_j as the **average** effect on Y of a one unit increase in X_j , **holding all other predictors fixed**. In the advertising example, the model becomes:

$$\text{sales} = \beta_0 + \beta_1 \times \text{TV} + \beta_2 \times \text{radio} + \beta_3 \times \text{newspaper} + \epsilon$$

Interpreting Regression Coefficients

- ▶ The ideal scenario is when the predictors are uncorrelated – a **balanced design**:
 - ▶ Each coefficient can be estimated and tested separately.
 - ▶ Interpretations such as “*a unit change in X_j is associated with a β_j change in Y , while all other variables stay fixed*”, are possible.
- ▶ Correlations amongst predictors cause problems:
 - ▶ The variance of all coefficients tends to increase, sometimes dramatically.
 - ▶ Interpretations become hazardous: when X_j changes, everything else changes.
- ▶ **Claims of causality** should be avoided for observational data.

The Woes of (Interpreting) Regression Coefficients

“Data Analysis and Regression” – Mosteller and Tukey (1977)

A regression coefficient β_j estimates the expected change in Y per unit change in X_j , **with all other predictors held fixed**. But predictors usually change together!

The Woes of (Interpreting) Regression Coefficients

“Data Analysis and Regression” – Mosteller and Tukey (1977)

A regression coefficient β_j estimates the expected change in Y per unit change in X_j , **with all other predictors held fixed**. But predictors usually change together!

- ▶ **Example 1:** Y = total change in pocket; X_1 = # of coins; X_2 = # of pennies, nickels, dimes. By itself, coefficient of Y on X_2 is > 0 . But with X_1 in the model?

The Woes of (Interpreting) Regression Coefficients

“Data Analysis and Regression” – Mosteller and Tukey (1977)

A regression coefficient β_j estimates the expected change in Y per unit change in X_j , **with all other predictors held fixed**. But predictors usually change together!

- ▶ **Example 1:** Y = total change in pocket; X_1 = # of coins; X_2 = # of pennies, nickels, dimes. By itself, coefficient of Y on X_2 is > 0 . But with X_1 in the model?
- ▶ **Example 2:** Y = number of tackles by a football player in a season; W and H are his weight and height. Fitted regression model is $\hat{Y} = b_0 + 0,50W - 0,10H$. How do we interpret $\hat{\beta}_2 < 0$?

Two Quotes by Famous Statisticians

George Box

“Essentially, all models are wrong, but some are useful.”

Two Quotes by Famous Statisticians

George Box

“Essentially, all models are wrong, but some are useful.”

Fred Mosteller and John Tukey (paraphrasing George Box)

“The only way to find out what will happen when a complex system is disturbed is to disturb the system, not merely to observe it passively.”

Estimation and Prediction for Multiple Regression


- ▶ Given estimates $\hat{\beta}_0, \hat{\beta}_1, \dots, \hat{\beta}_p$, we can make prediction using the formula:

$$\hat{y} = \hat{\beta}_0 + \hat{\beta}_1 x_1 + \hat{\beta}_2 x_2 + \dots + \hat{\beta}_p x_p$$

- ▶ We estimate $\hat{\beta}_0, \hat{\beta}_1, \dots, \hat{\beta}_p$ as the values that minimize the sum of squared residuals:

$$RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \sum_{i=1}^n (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_{i1} - \dots - \hat{\beta}_p x_{ip})^2$$

This is done using standard statistical software. The values $\hat{\beta}_0, \hat{\beta}_1, \dots, \hat{\beta}_p$ that minimize RSS are the multiple least squares regression estimates.



Results for Advertising Data

	Coefficient	Std. Error	<i>t</i> -statistic	<i>p</i> -value
Intercept	2.939	0.3119	9.42	<0.0001
TV	0.046	0.0014	32.81	<0.0001
radio	0.189	0.0086	21.89	<0.0001
newspaper	-0.001	0.0059	-0.18	0.8599

Correlations:

	TV	radio	newspaper	sales
TV	1.0000	0.0548	0.0567	0.7822
radio		1.0000	0.3541	0.5762
newspaper			1.0000	0.2283
sales				1.0000

Some Important Questions

1. Is at least one of the predictors $X_1, X_2, \dots, X_p$ useful in predicting the response?

Some Important Questions

1. Is at least one of the predictors $X_1, X_2, \dots, X_p$ useful in predicting the response?
2. Do all the predictors help to explain Y , or is only a subset of the predictors useful?

Some Important Questions

1. Is at least one of the predictors $X_1, X_2, \dots, X_p$ useful in predicting the response?
2. Do all the predictors help to explain Y , or is only a subset of the predictors useful?
3. How well does the model fit the data?

Some Important Questions

1. Is at least one of the predictors $X_1, X_2, \dots, X_p$ useful in predicting the response?
2. Do all the predictors help to explain Y , or is only a subset of the predictors useful?
3. How well does the model fit the data?
4. Given a set of predictor values, what response value should we predict, and how accurate is our prediction?

Some Important Questions

1. Is at least one of the predictors $X_1, X_2, \dots, X_p$ useful in predicting the response?
2. Do all the predictors help to explain Y , or is only a subset of the predictors useful?
3. How well does the model fit the data?
4. Given a set of predictor values, what response value should we predict, and how accurate is our prediction?

Is at Least One Predictor Useful?

For the first question, we can use the F -statistic:

$$F = \frac{\frac{(TSS-RSS)}{p}}{\frac{RSS}{n-p-1}} \sim F_{p, n-p-1}$$

Quantity	Value
Residual Standard Error	1.69
R^2	0.897
F -statistic	570

Deciding on the Important Variables

- ▶ **All Subsets / Best Subsets Regression:** Compute least squares fit for all possible subsets and choose based on a criterion balancing training error and model size.

Deciding on the Important Variables

- ▶ **All Subsets / Best Subsets Regression:** Compute least squares fit for all possible subsets and choose based on a criterion balancing training error and model size.
- ▶ However, there are 2^p possible models (e.g., $p = 40 \implies > 1$ billion models!).
- ▶ Instead we need an automated approach that searches through a subset of them. The book discusses two commonly use approaches next.

Variable Selection Methods

Forward Selection

- ▶ Begin with the **null model** (a model that contains an intercept but no predictors).
- ▶ Fit p simple linear regression and add to the null model the variable that results in the lowest RSS.
- ▶ Add to that model the variable that results in the lowest RSS amongst all two-variable models.
- ▶ Continue adding variables one by one until a stopping rule is satisfied. i.e. when all remaining variables have a *p-value* above some threshold.

Variable Selection Methods

Backward Selection

- ▶ Start with all variables in the model.
- ▶ Remove the variable with the largest p -value (that is, the variable that is the least statistically significant).
- ▶ The new $(p - 1)$ -variable model is fit, and the variable with the largest p -value is removed.
- ▶ Continue until a stopping rule is reached. i.e. we may stop when all remaining variable have a significant p -value defined by some significance threshold.

Model Selection Criteria

Other systematic criteria for choosing an optimal model path include:

- ▶ Mallow's C_p
- ▶ Akaike Information Criterion (AIC)
- ▶ Bayesian Information Criterion (BIC)
- ▶ Adjusted R^2
- ▶ Cross-Validation (CV)

Plan

Simple Linear Regression

- Estimating the Coefficients

- Assessing the Accuracy of the Coefficient Estimates

- Assessing the Accuracy of the Model

Multiple Linear Regression

- Estimating the Regression Coefficients

- Some Important Questions

Other Considerations in the Regression Model

- Qualitative Predictors


- Extensions of the Linear Model

- Potential Problems

Qualitative Predictors

- ▶ Qualitative predictors take a discrete set of values (also known as **categorical** predictors or **factor variables**).
- ▶ i.e. In addition to the 7 quantitative variables shown, there are four qualitative variables: **gender**, **student** (student status), **status** (marital status), and **ethnicity** (Caucasian, African American or Asian).

Credit card data



Regression

Qualitative Predictors –continued

- ▶ Example: investigate differences in credit card balance between males and females, ignoring the other variables.
- ▶ We create a new variable (a binary predictor, e.g. gender: Female/Male):

$$x_i = \begin{cases} 1 & \text{if } i\text{-th person is female} \\ 0 & \text{if } i\text{-th person is male} \end{cases}$$

- ▶ Resulting model:

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i = \begin{cases} \beta_0 + \beta_1 + \epsilon_i & \text{if female} \\ \beta_0 + \epsilon_i & \text{if male} \end{cases}$$

Qualitative Predictors –continued

- ▶ Example: investigate differences in credit card balance between males and females, ignoring the other variables.
- ▶ We create a new variable (a binary predictor, e.g. gender: Female/Male):

$$x_i = \begin{cases} 1 & \text{if } i\text{-th person is female} \\ 0 & \text{if } i\text{-th person is male} \end{cases}$$

- ▶ Resulting model:

$$y_i = \beta_0 + \beta_1 x_i + \epsilon_i = \begin{cases} \beta_0 + \beta_1 + \epsilon_i & \text{if female} \\ \beta_0 + \epsilon_i & \text{if male} \end{cases}$$

- ▶ Interpretation?

Credit card data –continued

Results for gender model:

	Coefficient	Std. Error	<i>t</i> -statistic	<i>p</i> -value
Intercept	509.80	33.13	15.389	<0.0001
gender[Female]	19.73	46.05	0.429	0.6690

Qualitative Predictors with More Than Two Levels

- ▶ With $k > 2$ levels, create $k - 1$ dummy variables.
- ▶ i.e. for **ethnicity** (Asian, Caucasian, African American) create two dummy variables:

$$x_{i1} = \begin{cases} 1 & \text{if Asian} \\ 0 & \text{otherwise} \end{cases}, \quad x_{i2} = \begin{cases} 1 & \text{if Caucasian} \\ 0 & \text{otherwise} \end{cases}$$

- ▶ Resulting model:

$$y_i = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \epsilon_i = \begin{cases} \beta_0 + \beta_1 + \epsilon_i & \text{if Asian} \\ \beta_0 + \beta_2 + \epsilon_i & \text{if Caucasian} \\ \beta_0 + \epsilon_i & \text{if African American (baseline)} \end{cases}$$

Qualitative Predictors with More Than Two Levels –continued

- Then both of these variables can be used in the regression equation, in order to obtain the model:

$$y_i = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \epsilon_i = \begin{cases} \beta_0 + \beta_1 + \epsilon_i & \text{if Asian} \\ \beta_0 + \beta_2 + \epsilon_i & \text{if Caucasian} \\ \beta_0 + \epsilon_i & \text{if African American (baseline)} \end{cases}$$

Note: There will always be one fewer dummy variable than the number of levels. The level with no dummy variable –AA in this example– is known as the **baseline**

Results for ethnicity

	Coefficient	Std. Error	<i>t</i> -statistic	<i>p</i> -value
Intercept	531.00	46.32	11.464	<0.0001
ethnicity[Asian]	-18.69	65.02	-0.287	0.7740
ethnicity[Caucasian]	-12.50	56.68	-0.221	0.8260

Extensions of the Linear Model

Removing the additive assumption: *interactions* and *nonlinearity*...

► Interactions:

- In our previous analysis of the **Advertising** data, we assumed that the effect on **sales** of increasing one advertising medium is independent of the amount spent on the other media.
- i.e. the linear model:

$$\text{sales} = \beta_0 + \beta_1 \times \text{TV} + \beta_2 \times \text{radio} + \beta_3 \times \text{newspaper}$$


states that the average effect on **sales** of a one-unit increase in TV is always β_1 , regardless of the amount spent on **radio**.

Extensions of the Linear Model –continued

► Interactions:

- But suppose that spending money on radio advertising actually increases the effectiveness of TV advertising, so that the slope term for TV should increase as **radio** increases.
- In this situation, given a fixed budget of \$100,000, spending half on **radio** and half on TV may increase **sales** more than allocating the entire amount to either TV or to **radio**.
- In marketing, this is known as a **synergy** effect, and in statistics it is referred to as an **interaction** effect.

Interaction in the Advertising data?



- ▶ When levels of either TV or radio are low, then the true **sales** are lower than predicted by the linear model.
- ▶ But when advertising is split between the two media, then the model tends to underestimate **sales**.

Modelling Interactions – Advertising Data Results

Model takes the form:

$$\begin{aligned}\text{sales} &= \beta_0 + \beta_1 \times \text{TV} + \beta_2 \times \text{radio} + \beta_3 \times (\text{TV} \times \text{radio}) + \epsilon \\ &= \beta_0 + (\beta_1 + \beta_3 \times \text{radio}) \times \text{TV} + \beta_2 \times \text{radio} + \epsilon\end{aligned}$$

	Coefficient	Std. Error	<i>t</i> -statistic	<i>p</i> -value
Intercept	6.7502	0.2480	27.23	< 0,0001
TV	0.0191	0.0015	12.70	< 0,0001
radio	0.0289	0.0089	3.24	0.0014
TV×radio	0.0011	0.0001	20.73	< 0,0001

Interpretation

- ▶ The results in this table suggests that interactions are important.
- ▶ The *p-value* for the interaction term $\text{TV} \times \text{radio}$ is extremely low, indicating that there is strong evidence for $H_A : \beta_3 \neq 0$.
- ▶ R^2 increases from 89.7 % (additive) to 96.8 % (with interaction).

Interpretation –continued

- ▶ Explains variability remaining after fitting additive model!
- ▶ This means that

$$\frac{96,8 - 89,7}{100 - 89,7} = 69 \%$$

of the variability in **sales** that remains after fitting the additive model has been explained by the interaction term.

- ▶ The coefficient estimates in the table suggest that an increase in TV advertising of \$1,000 is associated with increased sales of

$$(\hat{\beta}_1 + \hat{\beta}_3 \times \mathbf{radio}) \times 1000 = 19 + 1,1 \times \mathbf{radio} \text{ units.}$$

- ▶ An increase in radio advertising of \$1,000 will be associated with an increase in sales of

$$(\hat{\beta}_1 + \hat{\beta}_3 \times \mathbf{TV}) \times 1000 = 29 + 1,1 \times \mathbf{TV} \text{ units.}$$

The Hierarchy Principle

- Sometimes it is the case that an interaction term has a very small *p-value*, but the associated main effects (in this case, **TV** and **radio**) do not.

Hierarchy Principle

*If we include an interaction in a model, we should also include the main effects, even if the *p-values* associated with their coefficients are not significant.*

The Hierarchy Principle –continued

- ▶ The rationale for this principle is that interactions are hard to interpret in a model without main effects –their meaning is changed.
- ▶ Specifically, the interaction terms also contain main effects, if the model has no main effect terms.

Interactions between qualitative and quantitative variables

- ▶ Consider the **Credit** data set, and suppose that we wish to predict **balance** using **income** (quantitative) and **student** (qualitative).
- ▶ Without an interaction term, the model takes the form:



$$\begin{aligned}\text{balance}_i &\approx \beta_0 + \beta_1 \times \text{income}_i + \begin{cases} \beta_2 & \text{if } i\text{-th person is a student} \\ 0 & \text{if } i\text{-th person is not a student} \end{cases} \\ &= \beta_{\times} \text{income}_i + \begin{cases} \beta_0 + \beta_2 & \text{if } i\text{-th person is a student} \\ \beta_0 & \text{if } i\text{-th person is not a student} \end{cases}\end{aligned}$$

Interactions between qualitative and quantitative variables –continued

- With interactions, it takes the form:

$$\begin{aligned}\text{balance}_i &\approx \beta_0 + \beta_1 \times \text{income}_i + \begin{cases} \beta_2 + \beta_3 \times \text{income}_i & \text{if student} \\ 0 & \text{if not student} \end{cases} \\ &= \begin{cases} (\beta_0 + \beta_2) + (\beta_1 + \beta_3) \times \text{income}_i & \text{if student} \\ \beta_0 + \beta_1 \times \text{income}_i & \text{if not student} \end{cases}\end{aligned}$$


Interactions between qualitative and quantitative variables –continued



- Credit data; Left: no interaction between `income` and `student`. Right: with an interaction term between `income` and `student`.

Non-linear Effects of Predictors

polynomial regression on Auto data



Non-linear Effects of Predictors –continued

- ▶ Polynomial Regression extends linear models to non-linear relationships.
- ▶ The figure suggests that:

$$\text{mpg} = \beta_0 + \beta_1 \times \text{horsepower} + \beta_2 \times \text{horsepower}^2 + \epsilon$$

	Coefficient	Std. Error	<i>t</i> -statistic	<i>p</i> -value
Intercept	56.9001	1.8004	31.6	<0.0001
horsepower	-0.4662	0.0311	-15.0	<0.0001
horsepower ²	0.0012	0.0001	10.1	<0.0001

What We Did Not Cover & Generalizations

Potential Issues in Regression (Section 3.3.3)

Outliers, Non-constant variance of error terms, High leverage points, Collinearity.


Generalizations of the Linear Model

- ▶ **Classification:** Logistic regression, Support Vector Machines.
- ▶ **Non-linearity:** Kernel smoothing, Splines, GAMs, Nearest Neighbors.
- ▶ **Interactions:** Tree-based methods, Bagging, Random Forests, Boosting (these also capture non-linearities).
- ▶ **Regularized fitting:** Ridge regression, Lasso

End of Chapter 3.

An Introduction to Statistical Learning with Applications in R

a.k.a. ISLR



Content has been extracted from *An Introduction to Statistical Learning with Applications in R*, Second Edition, by James, Witten, Hastie & Tibshirani. Springer. 2023.

Visit <https://www.statlearning.com/>.